

Interactive - 08 - create unique id and a primary key

It is really hard to update data when you don't have a unique way of identifying a row. People will often imagine that there is something unique in the data. Some combination of name/date/place-of-birth, generally it is not true.

The easy way to deal with this is to create unique keys. PostgreSQL offers two ways of doing this. The first is a key that counts up. This makes short keys, that's good but also creates a bottle neck with a single source of "count-up" values. The second way is use a universal unique identifier or UUID. UUIDs can be created in parallel and will not collide. This has lots of advantages, but they are long, hard to type and take up lots of space in your database.

We will get back to sequential values later, for a different purpose. Let's do the more general approach and use UUIDs.

We are going to change our table to have a unique ID column, then create an index on that column.

This relies on having a way to generate UUIDs, it is an extension package to the default PostgreSQL database. (double quotes, " are correct for this command).

```
CREATE EXTENSION IF NOT EXISTS "uuid-oss";
```

There are ways to add a column to the table without re-creating the table. They will not work for this.

Let's save our data for later, then recreate the table.

First we are going to rename the table, so we can save it for later. (Familiar, we did this last time)

```
drop table if exists old_name_list;
ALTER TABLE name_list
    RENAME TO old_name_list;
```

Now we will re-create the table with the check constraint.

```
create table name_list (
    name_list_id UUID NOT NULL DEFAULT uuid_generate_v4() primary key,
    real_name text,
    age int check ( age > 0 and age < 154 ),
    state char varying (2)
);
```

Now copy the data back.

```
INSERT INTO name_list ( real_name, age, state )
    SELECT real_name, age, state
    FROM old_name_list;
```

Verify data amount. Check how many rows.

```
select count(1) from name_list;
select count(1) from old_name_list;
```

Provided that the counts are the same let's clean up the temporary table with the old data.

```
drop table old_name_list;
```

There are 2 parts that we added. Both a ID column that is unique and a thing called a PRIMARY KEY.

Your result should be similar to this list (the name_list_id's will be unique to you and will not match the UUIDs below).

name_list_id	real_name	age	state
1fb069ca-a4a2-412a-8401-c835b6c45cb7	Jane True	20	WY
39fff492-f369-4552-bf49-b8d7ae57bf9a	Tom Ace	31	NJ
8e91d65f-ee0a-4642-8d62-1cc1e1058238	Steve Pen	33	NJ
e245b3ec-9e5f-4112-ac7e-1d9732d9a4f1	Laura Jean Alkinoos	34	PA
b8c2468e-694d-46de-acc3-e8da239aa1a3	Philip Schlump	62	WY
983b17d2-39f7-45bc-96b2-66ffdbb96c0e	Bob True	22	WY
2db6af28-47c7-4107-9b67-bdec3321f14c	Liz Trubune	30	WY
ea51c6b2-37f0-47ed-9211-d1064d3eacfd	Jane True	44	WY
4c827537-1755-4242-8241-a77ff4173554	Lary Smith	58	NJ
3674b39d-4a61-45e0-99e6-65fb7edbbf82	Dave Dave	21	NJ
e543c622-24be-4cc2-8241-e8a2dfd1f15c	Laura Ann Alkinoos	34	PA

(11 rows)

The grading code will look at the table name_list and verify that you have 11 rows of data in it.

Tags: “primary key”, “uuid”, “unique id”, “UUID”

Validate: SQL-Select, “select ‘PASS’ as x from (select count(1) as x from name_list) as t1 where t1.x = 11”