

# ACID

---

In computer science, **ACID** (atomicity, consistency, isolation, durability) is a set of properties of database transactions intended to guarantee data validity despite errors, power failures, and other mishaps. In the context of databases, a sequence of database operations that satisfies the ACID properties (which can be perceived as a single logical operation on the data) is called a *transaction*. For example, a transfer of funds from one bank account to another, even involving multiple changes such as debiting one account and crediting another, is a single transaction.

In 1983,<sup>[1]</sup> Andreas Reuter and Theo Härder coined the acronym *ACID*, building on earlier work by Jim Gray<sup>[2]</sup> who named atomicity, consistency, and durability, but not isolation, when characterizing the transaction concept. These four properties are the major guarantees of the transaction paradigm, which has influenced many aspects of development in database systems.

According to Gray and Reuter, the IBM Information Management System supported ACID transactions as early as 1973 (although the acronym was created later).<sup>[3]</sup>

## Contents

---

### Characteristics

Atomicity

Consistency

Isolation

Durability

### Examples

Atomicity

Consistency failure

Isolation failure

Durability failure

### Implementation

Locking vs multiversioning

Distributed transactions

### See also

### References

## Characteristics

---

The characteristics of these four properties as defined by Reuter and Härder are as follows:

### Atomicity

Transactions are often composed of multiple statements. Atomicity guarantees that each transaction is treated as a single "unit", which either succeeds completely, or fails completely: if any of the statements constituting a transaction fails to complete, the entire transaction fails and the database is left unchanged. An atomic system must guarantee atomicity in each and every situation, including power failures, errors and crashes.<sup>[4]</sup> A guarantee of atomicity prevents updates to the database occurring only partially, which can cause greater problems than rejecting the whole series outright. As a consequence, the transaction cannot be observed to be in progress by another database client. At one moment in time, it has not yet happened, and at the next it has already occurred in whole (or nothing happened if the transaction was cancelled in progress).

An example of an atomic transaction is a monetary transfer from bank account A to account B. It consists of two operations, withdrawing the money from account A and saving it to account B. Performing these operations in an atomic transaction ensures that the database remains in a consistent state, that is, money is neither debited nor credited if either of those two operations fail.<sup>[5]</sup>

## Consistency

Consistency ensures that a transaction can only bring the database from one valid state to another, maintaining database invariants: any data written to the database must be valid according to all defined rules, including constraints, cascades, triggers, and any combination thereof. This prevents database corruption by an illegal transaction, but does not guarantee that a transaction is *correct*. Referential integrity guarantees the primary key – foreign key relationship.<sup>[6]</sup>

## Isolation

Transactions are often executed concurrently (e.g., multiple transactions reading and writing to a table at the same time). Isolation ensures that concurrent execution of transactions leaves the database in the same state that would have been obtained if the transactions were executed sequentially. Isolation is the main goal of concurrency control; depending on the method used, the effects of an incomplete transaction might not even be visible to other transactions.<sup>[7]</sup>

## Durability

Durability guarantees that once a transaction has been committed, it will remain committed even in the case of a system failure (e.g., power outage or crash). This usually means that completed transactions (or their effects) are recorded in non-volatile memory.

## Examples

---

The following examples further illustrate the ACID properties. In these examples, the database table has two columns, A and B. An integrity constraint requires that the value in A and the value in B must sum to 100. The following SQL code creates a table as described above:

```
CREATE TABLE acidtest (A INTEGER, B INTEGER, CHECK (A + B = 100));
```

## Atomicity

Atomicity is the guarantee that series of database operations in an atomic transaction will either all occur (a successful operation), or none will occur (an unsuccessful operation). The series of operations cannot be separated with only some of them being executed, which makes the series of operations "indivisible". A guarantee of atomicity prevents updates to the database occurring only partially, which can cause greater problems than rejecting the whole series outright. In other words, atomicity means indivisibility and irreducibility.<sup>[8]</sup> Alternatively, we may say that a logical transaction may be made of, or composed of, one or more (several), physical transactions. Unless and until all component physical transactions are executed, the Logical transaction will not have occurred – to the effects of the database. Say our Logical transaction consists of transferring funds from account A to account B. This logical transaction may be composed of several physical transactions consisting of first removing the amount from account A as a first physical transaction and then, as a second transaction, depositing said amount in account B. We would not want to see the amount removed from account A before we are sure it has been transferred into account B. Then, unless and until both transactions have happened and the amount has been transferred to account B, the transfer has not, to the effects of the database, occurred.

## Consistency failure

Consistency is a very general term, which demands that the data must meet all validation rules. In the previous example, the validation is a requirement that  $A + B = 100$ . All validation rules must be checked to ensure consistency. Assume that a transaction attempts to subtract 10 from A without altering B. Because consistency is checked after each transaction, it is known that  $A + B = 100$  before the transaction begins. If the transaction removes 10 from A successfully, atomicity will be achieved. However, a validation check will show that  $A + B = 90$ , which is inconsistent with the rules of the database. The entire transaction must be cancelled and the affected rows rolled back to their pre-transaction state. If there had been other constraints, triggers, or cascades, every single change operation would have been checked in the same way as above before the transaction was committed. Similar issues may arise with other constraints. We may have required the data types of both A and B to be integers. If we were then to enter, say, the value 13.5 for A, the transaction will be cancelled, or the system may give rise to an alert in the form of a trigger (if/when the trigger has been written to this effect). Another example would be with integrity constraints, which would not allow us to delete a row in one table whose primary key is referred to by at least one foreign key in other tables.

## Isolation failure

To demonstrate isolation, we assume two transactions execute at the same time, each attempting to modify the same data. One of the two must wait until the other completes in order to maintain isolation.

Consider two transactions:  $T_1$  transfers 10 from A to B.  $T_2$  transfers 20 from B to A.

Combined, there are four actions:

1.  $T_1$  subtracts 10 from A.
2.  $T_1$  adds 10 to B.
3.  $T_2$  subtracts 20 from B.
4.  $T_2$  adds 20 to A.

If these operations are performed in order, isolation is maintained, although  $T_2$  must wait. Consider what happens if  $T_1$  fails halfway through. The database eliminates  $T_1$ 's effects, and  $T_2$  sees only valid data.

By interleaving the transactions, the actual order of actions might be:

1.  $T_1$  subtracts 10 from A.
2.  $T_2$  subtracts 20 from B.
3.  $T_2$  adds 20 to A.
4.  $T_1$  adds 10 to B.

Again, consider what happens if  $T_1$  fails while modifying B in Step 4. By the time  $T_1$  fails,  $T_2$  has already modified A; it cannot be restored to the value it had before  $T_1$  without leaving an invalid database. This is known as a write-write failure, because two transactions attempted to write to the same data field. In a typical system, the problem would be resolved by reverting to the last known good state, canceling the failed transaction  $T_1$ , and restarting the interrupted transaction  $T_2$  from the good state.

## Durability failure

Consider a transaction that transfers 10 from A to B. First it removes 10 from A, then it adds 10 to B. At this point, the user is told the transaction was a success. However, the changes are still queued in the disk buffer waiting to be committed to disk. Power fails and the changes are lost, but the user assumes (understandably) that the changes persist.

## Implementation

---

Processing a transaction often requires a sequence of operations that is subject to failure for a number of reasons. For instance, the system may have no room left on its disk drives, or it may have used up its allocated CPU time. There are two popular families of techniques: write-ahead logging and shadow paging. In both cases, locks must be acquired on all information to be updated, and depending on the level of isolation, possibly on all data that may be read as well. In write ahead logging, durability is guaranteed by copying the original (unchanged) data to a log before changing the database. That allows the database to return to a consistent state in the event of a crash. In shadowing, updates are applied to a partial copy of the database, and the new copy is activated when the transaction commits.

## Locking vs multiversioning

Many databases rely upon locking to provide ACID capabilities. Locking means that the transaction marks the data that it accesses so that the DBMS knows not to allow other transactions to modify it until the first transaction succeeds or fails. The lock must always be acquired before processing data, including data that is read but not modified. Non-trivial transactions typically require a large number of locks, resulting in substantial overhead as well as blocking other transactions. For example, if user A is running a transaction that has to read a row of data that user B wants to modify, user B must wait until user A's transaction completes. Two phase locking is often applied to guarantee full isolation.

An alternative to locking is multiversion concurrency control, in which the database provides each reading transaction the prior, unmodified version of data that is being modified by another active transaction. This allows readers to operate without acquiring locks, i.e., writing transactions do not block reading transactions, and readers do not block writers. Going back to the example, when user A's transaction requests data that user B is modifying, the database provides A with the version of that data that existed when user B started his transaction. User A gets a consistent view of the database even if other users are changing data. One implementation, namely snapshot isolation, relaxes the isolation property.

## Distributed transactions

Guaranteeing ACID properties in a distributed transaction across a distributed database, where no single node is responsible for all data affecting a transaction, presents additional complications. Network connections might fail, or one node might successfully complete its part of the transaction and then be required to roll back its changes because of a failure on another node. The two-phase commit protocol (not to be confused with two-phase locking) provides atomicity for distributed transactions to ensure that each participant in the transaction agrees on whether the transaction should be committed or not.<sup>[9]</sup> Briefly, in the first phase, one node (the coordinator) interrogates the other nodes (the participants) and only when all reply that they are prepared does the coordinator, in the second phase, formalize the transaction.

## See also

---

- Eventual consistency (BASE)
- CAP theorem
- Concurrency control
- Java Transaction API
- Open Systems Interconnection
- Transactional NTFS
- Two-phase commit protocol
- CRUD

## References

---

1. Haerder, T.; Reuter, A. (1983). "Principles of transaction-oriented database recovery". *ACM Computing Surveys*. **15** (4): 287. CiteSeerX 10.1.1.115.8124 (<https://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.115.8124>). doi:10.1145/289.291 (<https://doi.org/10.1145%2F289.291>). S2CID 207235758 (<https://api.semanticscholar.org/CorpusID:207235758>).
2. Gray, Jim (September 1981). "The Transaction Concept: Virtues and Limitations" (<http://research.microsoft.com/~gray/papers/theTransactionConcept.pdf>) (PDF). *Proceedings of the 7th International Conference on Very Large Databases*. Cupertino, CA: Tandem Computers. pp. 144–154. Retrieved March 27, 2015.
3. Gray, Jim; Reuter, Andreas (1993). *Distributed Transaction Processing: Concepts and Techniques* (<https://archive.org/details/transactionproce0000gray>). Morgan Kaufmann. ISBN 1-55860-190-2.
4. "Atomic operation" ([http://www.webopedia.com/TERM/A/atomic\\_operation.html](http://www.webopedia.com/TERM/A/atomic_operation.html)). *webopedia.com*. Webopedia. Retrieved 2011-03-23. "An operation during which a processor can simultaneously read a location and write it in the same bus operation. This prevents any other processor or I/O device from writing or reading memory until the operation is complete."
5. Amsterdam, Jonathan. "Atomic File Transactions, Part 1" (<https://web.archive.org/web/20160303090517/http://archive.oreilly.com/pub/a/onjava/2001/11/07/atomic.html>). *O'Reilly*. Archived from the original (<http://archive.oreilly.com/pub/a/onjava/2001/11/07/atomic.html>) on 2016-03-03. Retrieved 2016-02-28.
6. C. J. Date, "SQL and Relational Theory: How to Write Accurate SQL Code 2nd edition", *O'reilly Media, Inc.*, 2012, pg. 180.
7. "Isolation Levels in the Database Engine", Technet, Microsoft, [https://technet.microsoft.com/en-us/library/ms189122\(v=SQL.105\).aspx](https://technet.microsoft.com/en-us/library/ms189122(v=SQL.105).aspx)
8. "Atomicity" ([https://docs.oracle.com/cd/E17276\\_01/html/programmer\\_reference/transapp\\_atomicity.html](https://docs.oracle.com/cd/E17276_01/html/programmer_reference/transapp_atomicity.html)). *docs.oracle.com*. Retrieved 2016-12-13.

9. Bernstein, Philip A.; Newcomer, Eric (2009). "Chapter 8". *Principles of Transaction Processing* (<http://web.archive.org/web/20100807151625/http://www.elsevierdirect.com/product.jsp?isbn=9781558606234>) (2nd ed.). Morgan Kaufmann (Elsevier). ISBN 978-1-55860-623-4. Archived from the original (<http://www.elsevierdirect.com/product.jsp?isbn=9781558606234>) on 2010-08-07.
- 

Retrieved from "<https://en.wikipedia.org/w/index.php?title=ACID&oldid=1000430149>"

---

**This page was last edited on 15 January 2021, at 02:12 (UTC).**

Text is available under the Creative Commons Attribution-ShareAlike License; additional terms may apply. By using this site, you agree to the Terms of Use and Privacy Policy. Wikipedia® is a registered trademark of the Wikimedia Foundation, Inc., a non-profit organization.